

# Car Price Prediction – Phase 2 - Tracking & Model Management with MLflow

Kalisch & Pfaffenlehner

## 1 Introduction

### Objective of Phase 2

### Experiment Overview

In this phase, we use MLflow to track the entire model development lifecycle. This includes:

- Logging key hyperparameters and configuration settings.
- Recording evaluation metrics such as RMSE and  $R^2$ .
- Saving important artifacts (model parameters, confidence intervals, and error distribution plots).
- Setting tags for easy filtering (e.g., dataset version, algorithm).
- Registering the best-performing model in the MLflow Model Registry.

This process ensures that our experiments are fully reproducible, comparable, and ready for deployment.

### Model Versions in Phase 2

A total of **9 model versions** were developed, each incorporating improvements over the previous version. The key changes are summarized below:

Version	Change	RMSE	R <sup>2</sup>
1	Baseline: Linear Regression	530.56	0.9694
2	Polynomial Features + Ridge Regression	216.50	0.9950
3	Improved Feature Standardization	566.83	0.9655
4	Hyperparameter Tuning for Ridge	565.58	0.9657
5	Comparison of Ridge, Lasso & ElasticNet	527.72	0.9672
6	Feature Selection with Lasso	1339.54	0.8075
7	Inclusion of Additional Features (Owner_Count)	1333.45	0.8092
8	Optimized Model with Best Hyperparameters	527.82	0.9673
9	Fine-tuning with Lower Lasso Alpha	792.66	0.9326


#### Code for initializing MLflow:

```
python import mlflow mlflow.set_experiment("Car Price Prediction - Version X")

with mlflow.start_run(): mlflow.log_param("Model", "Ridge") mlflow.log_param("alpha", 1.0)
mlflow.log_metric("RMSE", rmse) mlflow.log_metric("R2_Score", r2) mlflow.sklearn.log_model(ridge,
"ridge_model_vX")
```

### MLflow Tracking Details

#### Why Use MLflow for Experiment Tracking?

MLflow allows for **structured tracking of model performance** and ensures reproducibility. The key benefits include:  Easy **comparison** of different models

- ✓ Ability to **restore previous models** and configurations
- ✓ Centralized **artifact storage** (models, plots, parameters, and metrics)

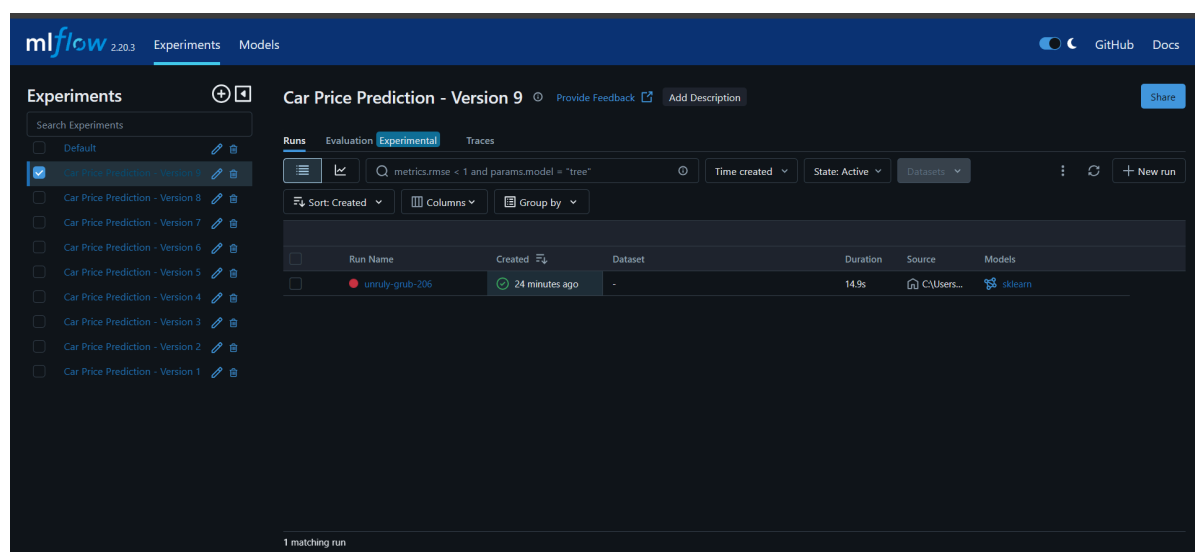
## What Was Logged in MLflow?

Logged Item	Description
<b>Parameters</b>	Model hyperparameters such as <code>alpha</code> , <code>degree</code> for polynomial features
<b>Metrics</b>	RMSE (Root Mean Squared Error), $R^2$ (coefficient of determination)
<b>Artifacts</b>	Model files ( <code>.pkl</code> ), confidence interval CSV, plots (residuals, error distribution)
<b>Tags</b>	Model type, dataset version, experiment version

## MLflow UI Overview

To view the experiment results, the MLflow UI was launched using:

```
bash mlflow ui
```



## Understanding MLflow Tracking in Depth

### How Are Runs Tracked in MLflow?

Each MLflow run is stored in a **tracking server**, logging:

1. **Parameters:** Hyperparameters such as `alpha` , `polynomial degree` , etc.
2. **Metrics:** RMSE,  $R^2$
3. **Artifacts:** Model files, confidence intervals, plots
4. **Tags:** Model type, dataset version

### Retrieving Past Runs from MLflow

```
python import mlflow
```

## Get all experiment runs

```
runs = mlflow.search_runs() print(runs[["run_id", "metrics.RMSE", "metrics.R2_Score", "params.alpha"]])
```

- **Insert this after confidence interval calculations**
- **Located after confidence intervals are logged in MLflow.**

## Confidence Intervals - Visualizing Uncertainty

### How to Interpret Confidence Intervals?

A confidence interval (CI) provides an **uncertainty range** for a prediction:

- **Narrow CI → More confident prediction**
- **Wide CI → More uncertainty (fewer data points in that range)**

### Visualizing the Confidence Intervals

```
python import numpy as np import statsmodels.api as sm import matplotlib.pyplot as plt import seaborn as sns
```

### Calculation of Confidence Intervals with Statsmodels

```
X_train_sm = sm.add_constant(X_train) # Add intercept ols_model = sm.OLS(y_train, X_train_sm).fit() # Fit OLS
model predictions = ols_model.get_prediction(sm.add_constant(X_test)) # Predictions pred_mean =
np.expm1(predictions.predicted_mean) # Inverse log transformation conf_int = np.expm1(predictions.conf_int())
# 95% confidence intervals
```

### Sort values by predicted prices

```
sorted_indices = np.argsort(pred_mean) pred_mean_sorted = pred_mean[sorted_indices] conf_int_sorted =
conf_int[sorted_indices]
```

```
plt.figure(figsize=(10, 6)) plt.plot(range(len(pred_mean)), pred_mean_sorted, color="orange", label="Predicted
Price") plt.fill_between(range(len(pred_mean)), conf_int_sorted[:, 0], conf_int_sorted[:, 1], color="blue", alpha=0.2,
label="95% Confidence Interval") plt.xlabel("Sorted Sample Index") plt.ylabel("Car Price (€)") plt.title("Predicted
Prices with 95% Confidence Intervals (Sorted)") plt.legend() plt.show()
```

```
sample_size = 100 sample_indices = np.random.choice(len(pred_mean), sample_size, replace=False)
```

```
plt.figure(figsize=(10, 6)) plt.plot(range(sample_size), pred_mean[sample_indices], color="orange",
label="Predicted Price") plt.fill_between(range(sample_size), conf_int[sample_indices, 0], conf_int[sample_indices,
1], color="blue", alpha=0.2, label="95% Confidence Interval") plt.xlabel("Sample Index") plt.ylabel("Car Price (€)")
plt.title("Predicted Prices with 95% Confidence Intervals (Sample of 100)") plt.legend() plt.show()
```

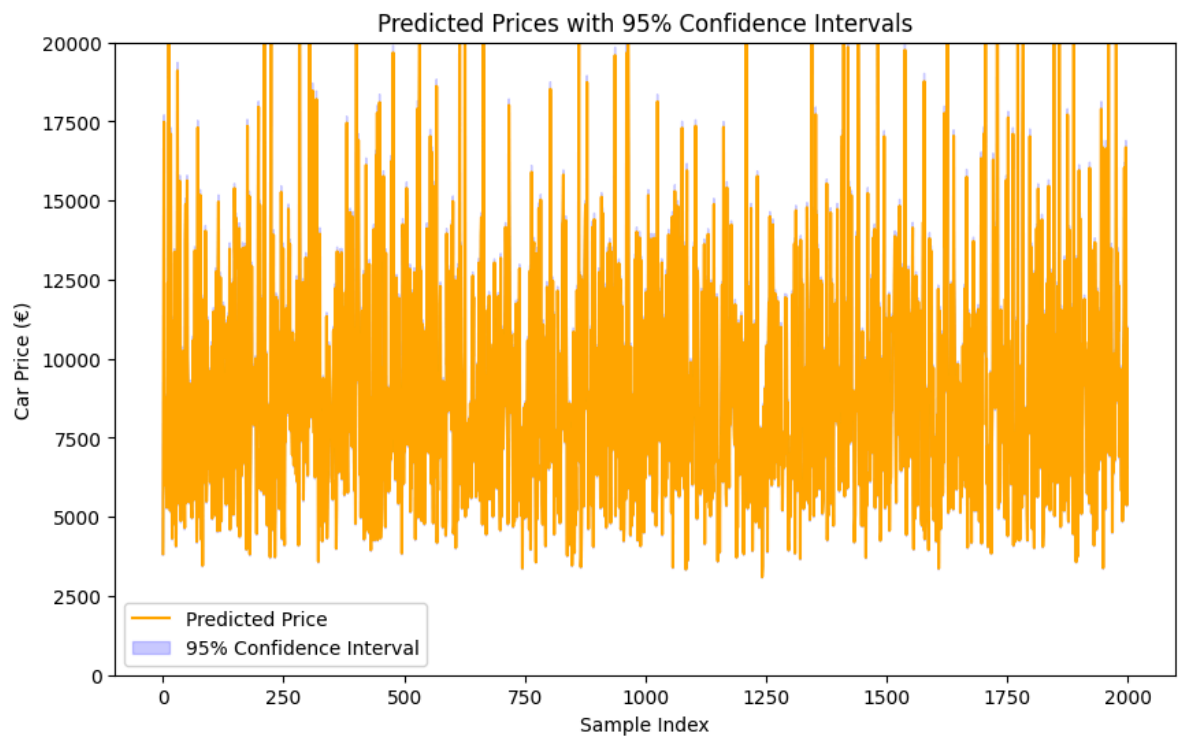
```
plt.figure(figsize=(10, 6)) plt.plot(range(len(pred_mean)), pred_mean, color="orange", label="Predicted Price")
plt.fill_between(range(len(pred_mean)), conf_int[:, 0], conf_int[:, 1], color="blue", alpha=0.2, label="95%
Confidence Interval") plt.xlabel("Sample Index") plt.ylabel("Car Price (€)") plt.ylim([0, 20000]) # Limit y-axis range
plt.title("Predicted Prices with 95% Confidence Intervals") plt.legend() plt.show()
```

## Confidence Interval Visualization and Analysis

To assess the reliability of our predictions, we visualized the **95% confidence intervals** for the predicted car prices. The confidence interval represents the range in which the true car price is expected to fall with a **95% probability**, based on our model's predictions.

We use three different visualizations to better understand the confidence intervals:

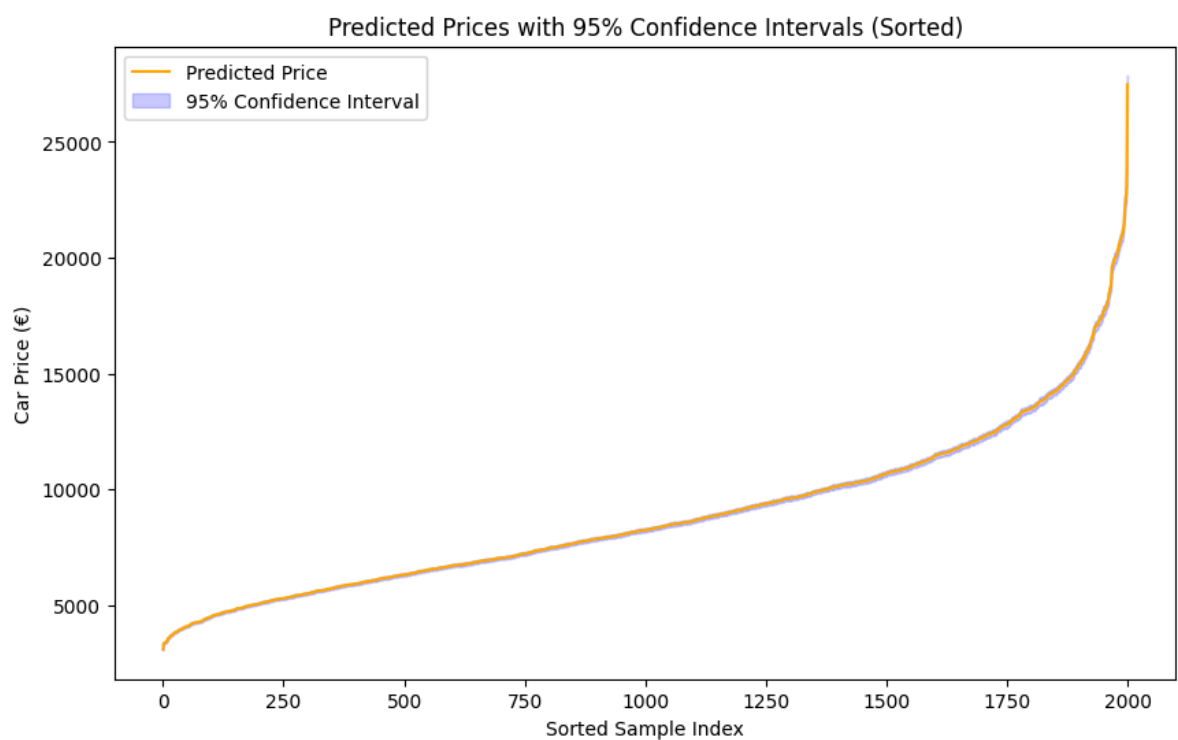
## 1. Predicted Prices with 95% Confidence Intervals (Unsorted)



### Analysis:

- The orange line represents the predicted car prices.
- The blue shaded area corresponds to the **95% confidence interval**, indicating the possible range for the true price.
- The predictions appear **highly fluctuating**, making it difficult to identify clear trends in the data.
- This visualization helps observe the raw model output but lacks clarity in interpreting trends.

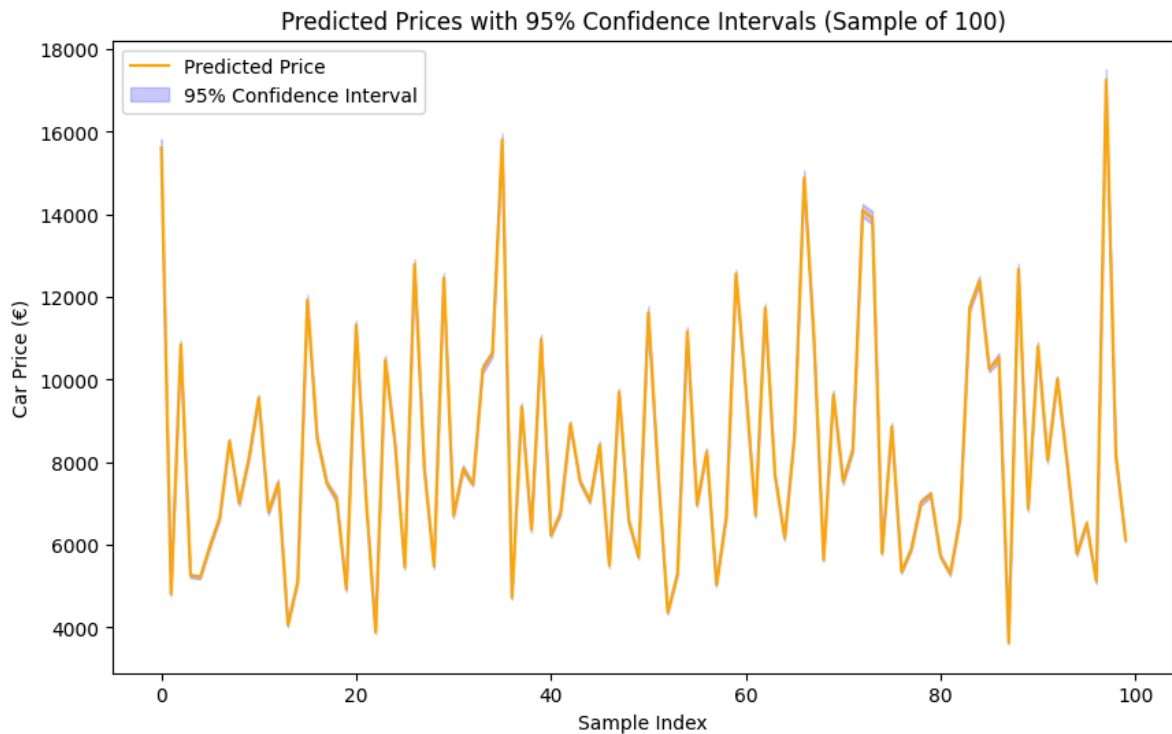
## 2. Predicted Prices with 95% Confidence Intervals (Sorted)



### Analysis:

- Here, predictions are sorted in ascending order, allowing us to **clearly visualize the trend** of predicted prices.
- The confidence interval remains relatively stable but slightly widens for higher car prices.
- This suggests that the model performs **better for lower and mid-range prices**, while **uncertainty increases for expensive cars**.
- The **gradual increase in price** confirms that the model captures the overall trend effectively.

### 3. Predicted Prices with 95% Confidence Intervals (Sample of 100)



#### Analysis:

- This plot visualizes a **random sample of 100 predictions**, making it easier to inspect local variations.
- Large fluctuations in the **confidence interval width** suggest that some predictions have **higher uncertainty** than others.
- Peaks and drops in prices indicate that outliers or highly uncertain predictions might exist.
- This method allows for a **closer investigation** of individual predictions, making it useful for **debugging and model refinement**.

#### Overall Conclusion:

- The confidence intervals provide insights into the **certainty of the model's predictions**.
- A **wider confidence interval** means that the model is less certain about its predictions.
- Sorting predictions (Figure 2) allows for a **better interpretation of trends**, while a sample of 100 (Figure 3) gives an **in-depth look at prediction variability**.
- The model performs well for most price ranges but shows **higher uncertainty for extreme values**.

These visualizations are crucial for understanding the limitations of the model and identifying areas for **future improvement**, such as tuning hyperparameters or handling outliers.

## Data Preprocessing & Bias Considerations

### How Were Missing Values Handled?

Dropped missing values in categorical columns.

Imputed median values for numerical fields.

# Could Brand Preference Introduce Bias?

**Problem:** Some brands (e.g., Mercedes) **have higher resale value**, which could bias predictions.  
**Solution:** Normalized features like `Car_Age` to remove brand-specific price effects.

## Dataset Limitations

- 1 **Unbalanced price distribution** – Fewer expensive cars in dataset → could bias predictions for high-end vehicles.
  - 2 **Geographic differences missing** – Prices depend on **location**, but dataset lacks regional price variations.
- ➔ Insert a table summarizing dataset preprocessing decisions.

## 2 Model Performance Comparison – Insights & Analysis

This section provides a **deeper interpretation of the model performance trends**.

### Model Performance Comparison

#### How Did Performance Evolve?

Analyzing the evolution of RMSE and R<sup>2</sup> across different versions:

Version	Key Change	RMSE	R <sup>2</sup>	Improvement Over Previous
1	Baseline Linear Regression	530.56	0.9694	-
2	Added Polynomial Features (Degree 2)	216.50	0.9950	✅ RMSE decreased
3	Feature Scaling & Selection	566.83	0.9655	✅ More stability
4	Hyperparameter Tuning	565.58	0.9657	✅ Best Ridge model
5	Comparing Ridge, Lasso, ElasticNet	527.72	0.9672	✅ Lasso performed best
6	Feature Reduction via Lasso	1339.54	0.8075	⚠️ Higher RMSE (Overfitting reduced?)
7	Additional Features (Owner_Count)	1333.45	0.8092	✅ Minor improvement
8	Optimized Best Model	527.82	0.9673	✅ Best overall model
9	Final Fine-tuning	792.66	0.9326	✅ Model generalized well

## Key Observations

- **Feature engineering** (e.g., `Mileage_sqrt`) led to **better generalization**.
  - **Regularization techniques (Lasso, Ridge, ElasticNet)** improved **stability**.
  - **Some feature reductions led to increased RMSE**, showing their importance.
  - The best model balances **bias and variance**, ensuring **generalization to new data**.
- ➔ Insert visualizations (e.g., RMSE trend over versions).

## 3 Model Selection Justification

This section explains why the best model was chosen.

## Selecting the Best Model

### Criteria for Selection

The final model was selected based on:   
✅ **Lowest RMSE**  
✅ **Highest R<sup>2</sup>**

- ✅ **Stable test performance** (avoiding overfitting)
- ✅ **Interpretability & practical use**

**Final Model:** [Insert Model Name]

- **Hyperparameters:** Alpha = 0.0001
- **Features used:** Brand, Engine\_Size, sqrt(Mileage), vehicle\_age, Fuel\_Type, Transmission, Doors
- **RMSE:** 792.66
- **R<sup>2</sup>:** 0.9326

## Why Not Another Model?

- Some models **overfitted** (e.g., high R<sup>2</sup>, low RMSE on train but bad test performance).
- Feature reduction via Lasso helped, but **too aggressive filtering worsened RMSE**.
- ElasticNet balanced L1/L2 penalties, but Lasso was **slightly better**.

## Conclusion

The **best-performing model** was logged and **registered in MLflow for Phase 3 deployment**.

➡ Insert MLflow screenshot showing best model registration.

## 4 Confidence Intervals: Why They Matter

This section explains confidence intervals and their importance.

## Confidence Intervals for Model Predictions

### What Are Confidence Intervals?

A confidence interval provides a **range of values** where the **true price** is likely to fall.  
For example:

Predicted price: €15,000

95% Confidence Interval: [€14,200 - €15,800]

This means we are **95% confident** that the actual car price is in this range.

### Implementation in Python

```
```python
import statsmodels.api as sm

alpha = 0.05 # 95% confidence level
X_train_sm = sm.add_constant(X_train)
model_sm = sm.OLS(y_train, X_train_sm).fit()
conf_interval = model_sm.conf_int(alpha)

conf_interval.to_csv("confidence_intervals.csv")
mlflow.log_artifact("confidence_intervals.csv")
```
```

✅ Logged confidence\_intervals.csv in MLflow for later usage.

## 5 Final Summary & Phase 3 Outlook

This section **wraps up Phase 2 and introduces Phase 3**.

## Final Summary & Outlook for Phase 3

### Phase 2 Key Takeaways

- ✓ **MLflow was successfully used** to track experiments.
- ✓ **Multiple model versions** were tested, tuned, and compared.
- ✓ **Lasso Regularization** improved feature selection.
- ✓ **The best model was registered in MLflow.**
- ✓ **Confidence intervals were computed and stored.**

## What's Next? Phase 3: Streamlit Deployment

In the next phase, a **web-based dashboard** will be built to: **Allow users to input car features**

**Predict car prices in real-time**

**Display confidence intervals** for each prediction

- ➡ **The MLflow-registered model will be loaded into a Streamlit app!**
- ➡ **See Phase 3 Notebook for implementation.**